

PROBLEM STATEMENT: To predict the rainfall based on various features of the dataset

```
In [1]: #importing libraries
import numpy as np
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv(r"C:\Users\user\Downloads\rainfall in india 1901-2015.csv")
df
```

Out[2]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan- Feb	Mar- May
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5	558.2	33.6	3373.2	136.3	560.3
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2	359.0	160.5	3520.7	159.8	458.3
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2	284.4	225.0	2957.4	156.7	236.1
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2	308.7	40.1	3079.6	24.1	506.9
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7	25.4	344.7	2566.7	1.3	309.7

Data cleaning and preprocessing

```
In [3]: df.head()
```

Out[3]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan- Feb	Mar- May	Jun- Sep
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5	558.2	33.6	3373.2	136.3	560.3	1696.3
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2	359.0	160.5	3520.7	159.8	458.3	2185.1
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2	284.4	225.0	2957.4	156.7	236.1	1874.1
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2	308.7	40.1	3079.6	24.1	506.9	1977.1
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7	25.4	344.7	2566.7	1.3	309.7	1624.1

In [4]:

df.tail()

Out[4]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	Mar-May	
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4	184.3	14.9	1533.7	7.9	196.2	10
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9	12.4	8.8	1405.5	19.3	99.6	11
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8	78.1	26.7	1426.3	60.6	131.1	10
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.2	59.0	62.3	1395.0	69.3	76.7	9
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4	231.0	159.0	1642.9	2.7	223.9	8

In [5]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4116 entries, 0 to 4115
Data columns (total 19 columns):
#   Column          Non-Null Count  Dtype
---  -
0   SUBDIVISION     4116 non-null   object
1   YEAR            4116 non-null   int64
2   JAN             4112 non-null   float64
3   FEB             4113 non-null   float64
4   MAR             4110 non-null   float64
5   APR             4112 non-null   float64
6   MAY             4113 non-null   float64
7   JUN             4111 non-null   float64
8   JUL             4109 non-null   float64
9   AUG             4112 non-null   float64
10  SEP             4110 non-null   float64
11  OCT             4109 non-null   float64
12  NOV             4105 non-null   float64
13  DEC             4106 non-null   float64
14  ANNUAL          4090 non-null   float64
15  Jan-Feb         4110 non-null   float64
16  Mar-May         4107 non-null   float64
17  Jun-Sep         4106 non-null   float64
18  Oct-Dec         4103 non-null   float64
dtypes: float64(17), int64(1), object(1)
memory usage: 611.1+ KB
```

In [6]:

df.shape

Out[6]:

(4116, 19)

In [7]:

df.describe()

Out[7]:

	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	
count	4116.000000	4112.000000	4113.000000	4110.000000	4112.000000	4113.000000	4111.000000	4109.000000	4112.000000	4110.000000
mean	1958.218659	18.957320	21.805325	27.359197	43.127432	85.745417	230.234444	347.214334	290.263497	197.361
std	33.140898	33.585371	35.909488	46.959424	67.831168	123.234904	234.710758	269.539667	188.770477	135.408
min	1901.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.400000	0.000000	0.000000	0.100
25%	1930.000000	0.600000	0.600000	1.000000	3.000000	8.600000	70.350000	175.600000	155.975000	100.525
50%	1958.000000	6.000000	6.700000	7.800000	15.700000	36.600000	138.700000	284.800000	259.400000	173.900
75%	1987.000000	22.200000	26.800000	31.300000	49.950000	97.200000	305.150000	418.400000	377.800000	265.800
max	2015.000000	583.700000	403.500000	605.600000	595.100000	1168.600000	1609.900000	2362.800000	1664.600000	1222.000

```
In [8]: df.isnull().sum()
```


```
Out[8]: SUBDIVISION    0  
YEAR                0  
JAN                  4  
FEB                  3  
MAR                  6  
APR                  4  
MAY                  3  
JUN                  5  
JUL                  7  
AUG                  4  
SEP                  6  
OCT                  7  
NOV                 11  
DEC                 10  
ANNUAL              26  
Jan-Feb             6  
Mar-May             9  
Jun-Sep            10  
Oct-Dec            13  
dtype: int64
```

```
In [9]: df.fillna(method="ffill",inplace=True)
```


```
In [10]: df['YEAR'].value_counts()
```

```
Out[10]: YEAR  
1963     36  
2002     36  
1976     36  
1975     36  
1974     36  
..  
1915     35  
1918     35  
1954     35  
1955     35  
1909     34  
Name: count, Length: 115, dtype: int64
```


```
In [11]: sns.lmplot(x='JUN',y='JUL',order=2,data=df,ci=None)  
plt.show()
```



```
In [12]: df=df[['JAN','FEB','MAR','APR','MAY']]  
sns.heatmap(df.corr(),annot=True)  
plt.show()
```



```
In [13]: sns.pairplot(df)
plt.show()
```




```
In [14]: x=np.array(df['FEB']).reshape(-1,1)
y=x=np.array(df['JAN']).reshape(-1,1)
```

```
In [15]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.30)
```


```
In [16]: regr=LinearRegression()
regr.fit(x_train,y_train)
print(regr.score(x_train,y_train))
```

```
1.0
```

```
In [17]: y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='blue')
plt.plot(x_test,y_pred,color='black')
plt.show()
```



```
In [18]: df500=df[:][:500]
sns.lmplot(x='MAY',y='MAR',order=2,ci=None,data=df500)
plt.show()
```




```
In [19]: df500.fillna(method='ffill',inplace=True)
```

```
In [20]: x=np.array(df500['FEB']).reshape(-1,1)
y=x*np.array(df500['JAN']).reshape(-1,1)
```

```
In [21]: X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.33)
         regr.fit(X_train,y_train)
         regr.fit(X_test,y_test)
```

```
Out[21]: ▾ LinearRegression
         LinearRegression()
```

```
In [22]: y_pred=regr.predict(X_test)
         plt.scatter(X_test,y_test,color='red')
         plt.plot(X_test,y_pred,color='purple')
         plt.show()
```



```
In [23]: from sklearn.metrics import r2_score
         model=LinearRegression()
         model.fit(X_train,y_train)
         y_pred=model.predict(X_test)
         r2=r2_score(y_test,y_pred)
         print("R2 Score:", r2)
```

R2 Score: 1.0

RIGDE MODEL

```
In [24]: from sklearn.linear_model import Lasso,Ridge
         from sklearn.preprocessing import StandardScaler
```

```
In [25]: features= df.columns[0:4]
         target= df.columns[-4]
```

```
In [26]: x=np.array(df['FEB']).reshape(-1,1)
         y=np.array(df['JAN']).reshape(-1,1)
```

```
In [27]: x= df[features].values
         y= df[target].values
         x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.4,random_state=100)
```

```
In [28]: ridgeReg=Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
train_score_ridge=ridgeReg.score(x_train,y_train)
test_score_ridge=ridgeReg.score(x_test,y_test)
```

```
In [29]: print("\n Ridge Model:\n")
print("the train score for ridge model is{}".format(train_score_ridge))
print("the test score for ridge model is{}".format(test_score_ridge))
```

Ridge Model:

```
the train score for ridge model is0.999999999828922
the test score for ridge model is0.999999999836248
```

```
In [30]: lr=LinearRegression()
```

LASSO MODEL


```
In [32]: print("\n Lasso Model:\n")
lasso=Lasso(alpha=10)
lasso.fit(x_train,y_train)
train_score_ls=lasso.score(x_train,y_train)
test_score_ls=lasso.score(x_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is{}".format(test_score_ls))
```

Lasso Model:

```
The train score for ls model is 0.999936226881545
The test score for ls model is0.9999362250950667
```

```
In [34]: pd.Series(lasso.coef_,features).sort_values(ascending=True).plot(kind="box")
```

Out[34]: <Axes: >




```
In [35]: from sklearn.linear_model import LassoCV
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,1,10],random_state=0).fit(x_train,y_train)
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))
```

```
0.9999999994453043
0.9999999993985869
```



```
In [37]: figure(figsize= (10,10))
plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker="*",markersize=5,color='maroon',label=r'Ridge;\$
plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'Ridge;\$alpha-grid$
plot(features,alpha=0.4,linestyle='none',marker='o',markersize=7,color="green",label='Linear Regression')
ticks(rotation = 90)
legend()
show()
```



ELASTIC NET

```
In [38]: from sklearn.linear_model import ElasticNet
         eln=ElasticNet()
         eln.fit(x,y)
         print(eln.coef_)
         print(eln.intercept_)
         print(eln.score(x,y))

[0.00000000e+00  9.99136156e-01  1.56645396e-04  0.00000000e+00]
0.01455738923123917
0.9999994172977724
```

```
In [39]: y_pred_elastic = eln.predict(x_train)
         mean_squared_error=np.mean((y_pred_elastic - y_train)**2)
         print(mean_squared_error)

0.000725960274965855
```

Conclusion: The above implemented models "Lasso and Ridge" regression is high accuracy compared to them

In []: